

PLACEMENT OS

Business Requirements Document

Placement Process Digitization for Premium B-Schools

Platform: iitiimcareers.in

Prepared for: Chittaranjan

Date: August 16, 2026

Status: Draft v2 — Pre-development requirements & design

Executive Summary

Placement OS replaces IIM Raipur CDPO's spreadsheet- and email-driven placement process — JD circulars, manual profile-packet compilation, ad-hoc shortlist tracking, and end-of-season reporting — with one connected pipeline, generalizable to other premium B-Schools. This v3 pass cross-checks the v2 requirements and data model directly against the CDPO's own working files (Students' Profile Sheet, Final Placement Datasheet, JPC Mailing Sheet) and adds the sections a build-ready BRD needs but v2 did not yet have: verified data mapping, success metrics, risks, privacy/compliance, and MVP acceptance criteria. Section and requirement numbering from v2 is preserved unchanged; all additions are appended as Appendices B–G so existing FR-ID references stay stable.

What's new in this pass:

- Appendix B — Data Validation & Migration Mapping: real column-level mapping from the three source files to the Section 5 data model, with concrete data-quality issues found (pivoted/merged layout in the Datasheet, TBD placeholders, a #REF! error, multi-date strings) that affect the season-1 import approach.
- Appendix C — Success Metrics & KPIs for the MVP, anchored to real 2024-25 season numbers (247 companies, 337 placed, 42 yet-to-be-placed) as a before/after baseline.
- Appendix D — Risks & Assumptions the build plan is currently resting on.
- Appendix E — Security, Privacy & Compliance, including India's DPDP Act 2023 implications for student PII (phone, email, marks, defaults status).
- Appendix F — MVP Acceptance Criteria / Definition of Done, module by module.
- Appendix G — Glossary of CDPO-specific terms (SPC, JPC, GL, defaults, merge status, etc.) for readers outside the placement office.

Recommended read order for a first-time reviewer: this summary → Section 11.3 (where Placement OS can/can't win) → Section 12.2 (build sequencing) → Appendix F (what "done" looks like for the pilot).

Table of Contents

1. Overview & Objectives

Placement OS digitizes the end-to-end campus placement process for premium B-Schools, starting with IIM Raipur's CDPO, and is designed to generalize to other premium institutes over time. It replaces the current spreadsheet- and email-driven process (JD circulars, manually compiled profile sheets, ad-hoc shortlist tracking, and end-of-season reporting) with a single connected pipeline from company onboarding through final placement reporting.

The product is a standalone application on its own subdomain, positioned as a premium SaaS offering under the iitimcareers.in platform. It is visually aligned with the platform's design language but is architecturally independent of the platform's other products (Cursivo, Job Search OS).

1.1 Objectives

- Eliminate manual, error-prone circulation of JDs and compilation of candidate packets.
- Give recruiters a self-serve portal for JD posting and shortlisting, reducing CDPO coordination overhead.
- Give the Placement Committee (SPCs) a live, single source of truth for pipeline status across all companies.
- Automate eligibility filtering, including compliance/defaults status, before a student can apply.
- Automate generation of the season-end Final Placement Datasheet in the format the CDPO office already uses.
- Digitize the pre-season outreach effort (company BD) so relationship history and call outcomes are tracked centrally instead of living in individuals' notes.

1.2 Source Documents Referenced

- IIM Raipur Students' Profile Sheet — PGP Batch 2023-25
- PGP 2024-25 Defaults Tracker (Forms, GL & Seminars, Summits, Process Default, Total Defaults)
- IIM Raipur Final Placement Datasheet 2023-25
- Sample recruiter JD circular email (CDPO → batch mailer format)
- JPC Mailing Sheet 2024-25 — the CDPO's existing per-coordinator outreach tracking sheets

2. User Personas & Roles

Five base personas interact with the system. A single individual can hold more than one role concurrently (e.g. a student who is also an SPC and also does BD outreach) — see Section 7 for how roles combine.

Persona	Who	Primary Goal	Key Pain Point Today
CDPO Admin	Placement Committee / CDPO office staff	Run a smooth, on-time placement season with accurate reporting	Manual JD circulation, spreadsheet-based tracking, error-prone final reporting
SPC	Student volunteers coordinating part of the process	Keep every company's shortlist round moving without drops	No live visibility into which company hasn't responded; manual nudging
Recruiter	Company HR / TA contact	Post a JD and get a clean shortlist quickly	Waits on the CDPO office to compile and email profile+CV packets manually
Student	Batch candidate	See what they're eligible for and track their applications	Learns about JDs only via mass email; no self-serve status tracking
BD / Outreach ("JPC")	SPC-role members doing pre-season company outreach — historically called JPCs	Convert prospects into companies that come to campus	Relationship history and call notes live in personal sheets/notebooks

Note on JPC vs. SPC: in the CDPO's existing process, "JPC" and "SPC" are the same underlying role — every JPC is an SPC-role member doing outreach. The distinction is supervisory, not a separate base role: a "Senior SPC" (a custom role cloned from SPC, see Section 8) oversees a set of JPCs' assigned companies, reviews their outreach, and can reassign ownership. Section 4.8 and Section 5 reflect this owner/supervisor pair.

3. User Journeys

Narrative walk-throughs of each persona's end-to-end experience, in sequence, tying together the feature set described in Section 4.

3.1 Recruiter Journey — "Posting a role and getting a shortlist"

Step 1 — Signup & approval

A HR contact from MindSprint signs up on Placement OS with a work email. The account is created in Pending status. A CDPO Admin reviews and approves it — the recruiter receives an activation email.

Step 2 — JD creation

The recruiter logs in and creates a new JD: role title, grade, CTC breakdown, location(s), eligible branches/specializations, eligible batch, an "unplaced students only" toggle, and an apply-by deadline. They attach a JD PDF and hit Publish (or Save Draft if a CDPO approval gate is enabled).

Step 3 — Eligible list preview

Before publishing, the system shows an estimated eligible-student count based on the criteria entered, so the recruiter can sanity-check their filters before the notification goes out.

Step 4 — Applications come in

As students apply, the recruiter's applicant list fills in — sortable by CGPA, work-ex, specialization. Sensitive fields (phone, personal email) stay masked at this stage.

Step 5 — Shortlisting

The recruiter selects candidates via bulk actions (shortlist / reject / waitlist) and can label rounds (GD / PI / Final). On shortlisting a candidate, their contact fields unmask automatically and the student is notified.

Step 6 — Iterating through rounds

As rounds progress, the recruiter updates each candidate's status. Each update automatically notifies the SPC and the student — no manual email needed.

Step 7 — Close-out

Once final selections are made, the recruiter marks the JD closed. Those records feed directly into the season's Final Placement Datasheet.

3.2 Student Journey — "From notification to offer"

Step 1 — Notification

A student receives an email in the familiar CDPO mailer format: company, role, CTC, location, deadline, eligibility note — because the system already filtered them in as eligible (branch/specialization match, unplaced status, defaults clear).

Step 2 — One-click apply

Clicking "Apply Now" opens an in-app screen showing their current Profile Sheet and CV are already on file (no re-entry), plus a live eligibility check confirming they qualify. They confirm and the application is submitted instantly.

Step 3 — Tracking

Their "My Applications" view shows every company they've applied to and the live status: Applied → Under Review → Shortlisted → Interview → Selected / Rejected / Waitlisted.

Step 4 — Shortlist notification

If shortlisted, they're notified automatically with round details coordinated by their SPC (slot/room/link).

Step 5 — Own status only

Throughout, the student can see their own defaults count and their own application statuses — not peers' outcomes for the same JD.

Step 6 — Outcome

On selection, their record locks in as Placed; they stop receiving "unplaced only" JD notifications automatically.

3.3 SPC Journey — "Keeping every pipeline moving"

Step 1 — Dashboard view

An SPC logs in to see every active company's current round and shortlist count in one dashboard, with a staleness flag on any company that hasn't updated status in a configurable number of days.

Step 2 — Logistics

For each round, the SPC assigns interview slots/rooms/links to shortlisted students and sends reminder nudges.

Step 3 — Escalation

If a recruiter goes stale on a round, the SPC (or Admin) follows up directly — the system surfaces this automatically rather than the SPC needing to track it manually in a spreadsheet.

Step 4 — Reporting

SPCs have read-only access to the placement dashboard, so they can answer batch-mates' questions about overall progress without needing Admin to pull numbers.

3.4 CDPO Admin Journey — "Running the season"

Step 1 — Season setup

At the start of the season, Admin creates the batch (e.g. PGP 2024-26), bulk-imports the student roster from the existing Profile Sheet format, and imports defaults data from the existing Defaults Tracker sheets.

Step 2 — User approvals

As recruiters and SPCs sign up, Admin reviews and approves accounts.

Step 3 — Oversight

Admin has full visibility across every company's pipeline, can override a shortlist status if needed, and monitors the audit log for all JD approvals, shortlist overrides, and role changes.

Step 4 — Reporting

Throughout the season, Admin pulls the live placement dashboard (placed/unplaced, average/median/highest CTC, company-wise breakdown). At season end, Admin exports the Final Placement Datasheet in the existing format for official CDPO reporting.

Step 5 — Access wind-down

Student and SPC accounts auto-expire at season end / graduation. Recruiter and BD accounts persist for next season's outreach.

3.5 BD / Outreach Journey — "Bringing companies to campus"

This journey is built directly on the CDPO's existing JPC Mailing Sheet workflow, carried over module-for-module from OutreachOS (excluding Inbox and Auto-Apply — see Section 4.8.1 for the full mapping).

Step 1 — Assignment

Ahead of the season, a Senior SPC imports a target-company list into the Company Pipeline (Prospect → Contacted → Interested → Committed → Onboarded) and assigns each company to a JPC (an SPC-role member) as its Owner.

Step 2 — Persona selection

Before reaching out, the JPC picks the company-type persona that matches the prospect — IT/Product, BFSI, Consulting, Core/Manufacturing, FMCG & Sales, Startup, PSU/Government, or Analytics & Data — which loads that category's pitch template and tone.

Step 3 — Outreach send

The JPC sends (or mail-merges) the pitch under their own name/signature. The system tracks per-recipient status: Email Sent → Opened → Clicked → Responded, or Not Interested / Call Back Later / Bounced.

Step 4 — Call logging

After any call, the JPC logs Call Remarks against that company (what was discussed, next step) — no recording, notes only, matching current practice.

Step 5 — Supervision

The Senior SPC reviews progress across all their JPCs' assigned companies, can add SPC Remarks, and reassigns ownership if a company goes stale.

Step 6 — JD form received

When a company confirms it will recruit, the JPC or Senior SPC marks "JD Form Received" on that company — a standalone manual flag; someone then separately creates the actual JD in Company Onboarding (Section 4.1), the two are not auto-linked.

Step 7 — Conversion & handoff

Once formally onboarded, the company card moves to Onboarded — the company can now self-serve signup and post a JD, closing the loop into the Recruiter journey (Section 3.1).

Step 8 — Funnel reporting

Admin sees the outreach funnel — contacted → responded → onboarded — per JPC and per season.

4. Functional Requirements

Grouped by module. Each requirement is tagged with an ID for traceability into future build tickets.

4.1 Company Onboarding & JD Management

ID	Requirement
FR-1.1	Recruiters can self-register with a work email; account activates only after Admin approval.
FR-1.2	Recruiters can create a JD: title, grade, CTC (fixed/variable/total), location(s), eligible branch/specialization, eligible batch, min CGPA/backlog rule, "unplaced only" toggle, open positions count, JD attachment.
FR-1.3	JD lifecycle: Draft → (optional Admin approval) → Published → Applications Closed → Shortlisting → Closed.
FR-1.4	System shows an estimated eligible-student count before publish.
FR-1.5	Recruiters with prior seasons can view JD history and clone a past JD as a template.
FR-1.6	A company may have multiple recruiter logins with identical permissions, scoped to that company's own JDs only.

4.2 Eligibility Engine & Notification

ID	Requirement
FR-2.1	System auto-matches JD criteria (branch, specialization, CGPA, backlog, placement status, defaults status) against the live roster to build the eligible list.
FR-2.2	"Unplaced only" JDs automatically exclude already-placed students.
FR-2.3	Students exceeding the configured defaults threshold are automatically excluded, with the reason visible to the student on self-check.
FR-2.4	Notification email replicates the existing CDPO mailer format (company, role, CTC, location, deadline, eligibility note, Apply CTA).
FR-2.5	Admin can manually add/remove specific students from an eligible list before the notification sends.
FR-2.6	System logs notification delivery (sent, and open-tracking where available) per student per JD.

4.3 Student Application Flow

ID	Requirement
FR-3.1	Students can apply in one click; existing Profile Sheet + latest CV are auto-attached — no re-entry.
FR-3.2	Students can withdraw an application any time before the deadline.
FR-3.3	"My Applications" view shows live status per company: Applied → Under Review → Shortlisted → Interview → Selected / Rejected / Waitlisted.

ID	Requirement
FR-3.4	Students can self-check eligibility and see the specific reason if blocked (e.g. open defaults).

4.4 Candidate Packet & Recruiter Shortlisting

ID	Requirement
FR-4.1	On application, system auto-compiles the Profile Sheet (personal info, up to 3 prior employers, grad/PG/10th-12th, credentials) + latest CV into a single packet.
FR-4.2	Recruiter applicant list is sortable/filterable by CGPA, work-ex, specialization.
FR-4.3	Bulk shortlist / reject / waitlist actions, with round labels (Round 1 / GD / PI / Final).
FR-4.4	Recruiters can download all packets as a single merged file, or view inline.
FR-4.5	Recruiters can add private notes per candidate, not visible to the student.
FR-4.6	Shortlist submission auto-notifies the affected students and the assigned SPC.

4.5 Shortlist Coordination (SPC)

ID	Requirement
FR-5.1	SPC dashboard shows all active JDs with current round and shortlist count.
FR-5.2	SPCs can assign interview slots/rooms/links to shortlisted students and trigger reminder nudges.
FR-5.3	Two-way status sync: a recruiter round-outcome update auto-notifies the SPC and the student.
FR-5.4	System flags a company/JD as stale if no round update within a configurable number of days.

4.6 Defaults & Compliance Tracker

ID	Requirement
FR-6.1	Mirrors the existing Defaults Tracker structure: per-activity records (GLs, seminars, summits, SIP/process defaults) rolling up to a Total Defaults count per student.
FR-6.2	Admin can configure the defaults threshold that blocks/warns JD eligibility.
FR-6.3	Bulk CSV import of attendance/default data from external sheets/forms, mapped to activity columns.
FR-6.4	Students can view their own default count and which activities are pending, read-only.

4.7 Final Placement Reporting

ID	Requirement
FR-7.1	Live dashboard: placed/unplaced counts, company-wise breakdown, average/median/highest CTC.

ID	Requirement
FR-7.2	One-click export of the Final Placement Datasheet matching the existing wide-table format (company ↔ student ↔ package).
FR-7.3	Configurable export templates for different audiences (official CDPO report, public placement report, accreditation submission).
FR-7.4	Historical batch-over-batch comparison once multiple seasons exist on the platform.

4.8 Pre-Season Outreach CRM

Built by carrying over OutreachOS's module set, minus Inbox and Auto-Apply, and adapted from a job-seeker-outreach tool into a company-outreach tool. Validated directly against the CDPO's existing JPC Mailing Sheet (per-JPC tabs, contact fields, call/SPC remarks, JD-received flag, merge status).

4.8.1 OutreachOS → Placement OS Module Mapping

OutreachOS Module	Placement OS Module	Adaptation
Dashboard (Home)	BD Home	Funnel stats (contacted → responded → onboarded) per JPC and per season, in place of application stats.
Job Discovery	— Dropped	Not needed for MVP (per your call).
Job Board	— Dropped	Folded into Company Pipeline — companies are entered directly rather than scored from a pasted JD.
Job Tracker	Company Pipeline	Same Kanban mechanic (Prospect → Contacted → Interested → Committed → Onboarded); card carries Owner (JPC) + Supervisor (Senior SPC) instead of a job card's role/company fields.
Compose	Outreach Composer	Persona swaps from "target role" to "company-type" (8 categories, Section 4.8.2); sender is always the logged-in JPC/SPC, matching current practice.
Scheduled	Scheduled Sends	Unchanged mechanic — queue, retry, reschedule of outreach mail.
Verified Contacts / Directory	Company Contact Directory	Per-company POC records, replacing personal-contact verification with recruiter-contact records (Section 4.8.3).
Templates	Outreach Templates	One template set per company-type persona instead of per target-role persona.
Deliverability	Deliverability Health	Unchanged — SPF/DKIM/DMARC + bounce-rate trend monitoring for the institute's sending domain, critical at mail-merge scale.
Profile Vault	Placement Committee Vault	Batch profile deck, brochure, and past placement stats replace a personal

OutreachOS Module	Placement OS Module	Adaptation
		resume/work-ex vault — used to personalize and attach to outreach.
Inbox	— Excluded	Per your instruction.
Auto-Apply	— Excluded	Per your instruction.

4.8.2 Company-Type Personas

Eight fixed persona/template categories, each defining pitch tone and talking points. The sender identity is always the logged-in user, never the persona — only the template content changes.

- IT/Product
- BFSI
- Consulting
- Core/Manufacturing
- FMCG & Sales
- Startup
- PSU/Government
- Analytics & Data

4.8.3 Functional Requirements

ID	Requirement
FR-8.1	Company pipeline (Kanban: Prospect → Contacted → Interested → Committed → Onboarded), with target-company list import.
FR-8.2	Every company has an assigned Owner (JPC) and, via their reporting line, a Supervisor (Senior SPC).
FR-8.3	Contact directory per company: Title, Full Name, Last Name, HR Designation, Email, cc email, Phone — multiple POCs per company, matching the existing sheet's fields.
FR-8.4	Outreach Composer with company-type persona selection (Section 4.8.2) that loads the matching template; sender is always the logged-in JPC/SPC.
FR-8.5	Mail-merge send with per-recipient status tracking: Email Sent, Email Opened, Email Clicked, Responded, Not Interested, Call Back Later, Bounced.
FR-8.6	Scheduled/queued sends with retry and reschedule (mirrors OutreachOS Scheduled tab).
FR-8.7	Manual Call Remarks log per company: free text, timestamp, logged-by — no recording integration in this phase.
FR-8.8	Separate SPC Remarks field for the supervising Senior SPC, distinct from the JPC's Call Remarks.
FR-8.9	"JD Form Received" manual Yes/No flag per company — a standalone tracking flag, not auto-linked to Company Onboarding (Section 4.1).
FR-8.10	Senior SPC/Admin can reassign a company's Owner if it goes stale.

ID	Requirement
FR-8.11	Deliverability health monitoring (SPF/DKIM/DMARC, bounce-rate trend) for the sending domain.
FR-8.12	Placement Committee Vault: batch profile deck, brochure, past placement stats, referenceable/attachable in outreach sends.
FR-8.13	Outreach funnel reporting: contacted → responded → onboarded, per JPC and per season.

4.9 Platform / Admin

ID	Requirement
FR-9.1	Batch/cohort setup with bulk student roster import from the existing Profile Sheet format.
FR-9.2	Role-based access control per Section 7.
FR-9.3	Audit log of key actions: JD approvals, shortlist overrides, role/permission changes — viewable by all Admins, editable/deletable by none.
FR-9.4	Login via college/institute SSO.

4.10 Resume Maker for Students

New module, built on Cursivo's CV editor and AI-assist stack, embedded directly in the student placement workflow. The CV produced here is what auto-attaches to every application (Section 4.3) and feeds the recruiter candidate packet (Section 4.4).

4.10.1 Cursivo → Placement OS Module Mapping

Cursivo Module	Placement OS Module	Adaptation
Template Gallery + TplPlacementCell	CV Templates	Placement-Cell layout (Academic Record promoted above Experience) set as default; full template library available as alternates.
Editor Sections (Academics, Experience, Projects, Positions, Skills, Certifications, Awards)	CV Editor	Core section-by-section CV data entry, unchanged.
cvSlice / multiCvSlice	Multi-CV Management	Students maintain multiple tailored CV versions, one per company-type persona (Section 4.8.2).
ATSPanel / MultiATSPanel / JdCoverageBar / atsRules	JD-Fit & ATS Scorer	Score a student's CV against a specific recruiter's JD before applying.
resumeParse / pdfTextExtraction / docxBUILDER / useCvExport	Import & Export	Parse an existing CV in; export polished PDF/Word out for application packets.
placement.js	Placement Season Mode	Deadline-urgency indicators and daily edit streaks, tied to each company's apply-by date (Section 4.2).

Cursivo Module	Placement OS Module	Adaptation
OnboardingWizard / SmartPasteModal	Fast CV Setup	Guided first-CV creation pre-filled from the existing Profile Sheet data — no blank-page start.
SharedCVViewer	Shared CV Link	Powers the CV view inside the recruiter candidate packet (Section 4.4).
AlDiffPanel / FullRewritePanel / CoPilotPanel / AchievementWizard	AI Writing Assist	AI-assisted bullet rewriting and structured achievement drafting, included in MVP.
ReviewInboxPanel (adapted)	SPC Review & Comments	SPC leaves inline comments anchored to a specific bullet or section; student applies or dismisses each one.

4.10.2 Functional Requirements

ID	Requirement
FR-10.1	Students build and maintain multiple CV versions from a single profile, one per company-type persona.
FR-10.2	CV editor auto-populates from the existing Profile Sheet on first use — no blank-page start.
FR-10.3	Students can score a CV against a specific company's JD and see section-by-section coverage gaps.
FR-10.4	AI writing assist rewrites bullets and drafts achievement statements from three guided prompts (action / outcome / metric) — never invents facts the student hasn't supplied.
FR-10.5	SPCs can leave inline comments on any student CV, anchored to a specific bullet or section; students see and action each comment individually.
FR-10.6	The most recent CV version auto-attaches to every application (Section 4.3) and candidate packet (Section 4.4) — no manual re-upload.
FR-10.7	CVs export to PDF and Word, matching the format recruiters expect.
FR-10.8	Placement Season Mode surfaces per-company deadline urgency directly in the CV editor, sourced from each JD's apply-by date.

5. Data Model

Core entities, derived from the existing Profile Sheet, Defaults Tracker, and Final Placement Datasheet formats plus the new pipeline/CRM/access-control concepts.

Entity	Key Fields
Student	roll_no, name, personal_info, total_work_ex, prior_employers[1-3], graduation_details, pg_details, 10th_12th_details, credentials, other_qualifications, placement_status
Company	name, sector, pipeline_stage (Prospect/Contacted/Interested/Committed/Onboarded), owner_user_id (JPC), supervisor_user_id (Senior SPC), jd_form_received (bool), past_hiring_history
CompanyContact	company_id, title, full_name, last_name, hr_designation, email, cc_email, phone
OutreachActivity	company_id, contact_id, channel (call/email), call_remarks, spc_remarks, previous_mails_summary, merge_status (Sent/Opened/Clicked/Responded/Not Interested/Call Back Later/Bounced), logged_by, timestamp
CompanyTypePersona	id, category_name (IT/Product, BFSI, Consulting, Core/Manufacturing, FMCG & Sales, Startup, PSU/Government, Analytics & Data), template_content
JD	company_id, role, grade, ctc_breakdown, location[], eligibility_rules, apply_by_deadline, status
Application	student_id, jd_id, status, round_history[], timestamps
DefaultRecord	student_id, activity_name, activity_type (GL / Summit / Process / Seminar), attended (bool), category_total
PlacementRecord	student_id, company_id, final_ctc, role, offer_date (season-end)
User	id, name, email, sso_identity_id, status, batch_id (nullable), company_id (nullable)
Role	id, name, is_base_role, cloned_from_role_id, default_permission_sets[]
PermissionSet	id, name, description, actions[] (module:action pairs)
FieldVisibilityRule	entity, field, applies_to_permission_set, unlock_condition
AuditLogEntry	actor_user_id, action, target_entity, target_id, timestamp, metadata
CVDocument	student_id, persona_id (company-type), version_no, template_id, ats_score, jd_coverage_score, file_url, updated_at
CVReviewComment	cv_document_id, spc_user_id, anchor_section, anchor_bullet_id, comment_text, status (open/applied/dismissed), timestamp

6. UI Reference — Key Screens

Wireframe-level references for the primary screen in each journey. Final visual design follows the platform's dark, premium design language; recruiter/admin screens skew toward dense data tables, student-facing screens stay close to the mailer format students already recognize.

6.1 Recruiter — JD Creation

New Job Description Save Draft Publish

Role Title Bid Manager	Grade / Band F3 / F4	CTC (Fixed) 16.5 LPA	CTC (Variable) 2.24 LPA
---------------------------	-------------------------	-------------------------	----------------------------

Location
☒ Chennai ☒ Bangalore

Eligible Branch
☒ All ☐ Finance ☐ Marketing ☐ Operations

Eligible Batch
PGP 2024-26 ▾

Unplaced only
☒ Yes ☐ No

Apply-by
20 Nov 2026, 09:59 AM ▾

JD Attachment

Estimated eligible students: 214 Preview List

6.2 Student — Notification & Apply

M MindSprint is hiring at your campus

Role Bid Manager (Grade F3 / F4)	CTC INR 18.74 LPA (88% Fixed) + Insurance + Relocation
-------------------------------------	---

Location
Chennai or Bangalore

Apply by
09:59:59 AM, 20 Nov 2026

Unplaced students only

View JD Apply Now →

— after tapping Apply Now —

Applying to MindSprint — Bid Manager

Profile Sheet up to date	View
CV attached (Resume_v3.pdf)	Change
Eligibility check	Passed — 0 open defaults

Confirm Application

6.3 Recruiter — Applicant Shortlisting

MindSprint — Bid Manager

142 applicants

Round 1 ▾

	Name	CGPA	Work-Ex	Specialization	Status
☒	A. Sharma	3.7	14 mo	Finance	Applied
☒	R. Iyer	3.9	0	Marketing	Applied
☐	K. Verma	3.4	26 mo	Operations	Applied
☒	S. Nataraj	3.8	8 mo	Finance	Applied
☐	P. Bose	3.5	0	Marketing	Applied

⋮ 137 more applicants

Selected: 2

Download

Shortlist →

6.4 SPC — Coordination Dashboard

Active Companies

MindSprint	Round: PI	Shortlisted: 12	⚠ 2d stale
Deloitte USI	Round: GD	Shortlisted: 28	
Amazon	Round: Final Offer	Shortlisted: 4	

Upcoming: MindSprint PI slots — 6 unassigned

6.5 Admin — Placement Report Dashboard

PGP 2024-26 — Placement Season [Export Final Datasheet](#)

Placed
337 / 480

Avg CTC
₹ 14.2 L

Median CTC
₹ 12.8 L

Company-wise Breakdown [Table ▼](#)

Company	Students Placed	Avg CTC
Amazon	6	21.4 L
D.E. Shaw	3	28.0 L
Cognizant Consulting	9	11.2 L

6.6 BD/Outreach — Company Pipeline

Company Pipeline [Mail Merge →](#) [+ New Company](#)

Prospect

ABC Corp
DEF Inc

Contacted

XYZ Ltd
PQR Co

Interested

MindSprint

Committed

Amazon
Deloitte

MindSprint [Persona: BFSI ▼](#) **RESPONDED**

Owner (JPC)
Rahul Iyer

Supervisor
Priya Nair (Sr. SPC)

JD Form Received
☐ No

Contact
 **Ms. Aishwarya S, HR Manager**
aishwarya.s@mindsprint.com · +91 73497 25719

Call Remarks

12 Nov — Rahul
"Spoke to HR head, interested; follow up after budget approval"

+ Log Call

SPC Remarks

13 Nov — Priya
"Good lead, keep following up weekly"

+ Add Remark

6.7 Admin — User Management

User Management

+ Invite User

Pending Approval (3)

M

R. Mehta

MindSprint (Recruiter)

Approve

Reject

B

A. Bose

SPC — self-nominated

Approve

Reject

All Users

Filter: Role ▾

P

Priya Nair

Student, SPC, BD · Batch PGP 24-26

Edit

R

Rahul Iyer

Recruiter (MindSprint)

Edit

C

CDPO Office

Admin

Edit

6.8 Admin — Role & Permission Management

Roles

+ New Role

Admin, SPC, BD, Recruiter, Student

base roles · locked

Senior SPC

custom, cloned from SPC

Edit

New Role

Clone from

SPC ▾

Name

Senior SPC

Permission Sets Included

☒ Shortlist Oversight
 ☒ Student Data — Full
 ☐ Reports — View Only → ☒ Reports & Export (upgraded)
 ☐ + Add another Permission Set

Create Role

6.9 Student — Resume Builder & SPC Review

My CV — Resume_v3

IT/Product persona

Export DOCX

Export PDF

Template: Placement Cell v2

Change

JD-Fit vs Amazon SDE

78%

Academic Performance Record

IIM Raipur — PGP	2024-26	CGPA: 8.2	Edit
NIT Rourkela — B.Tech	2016-20	CGPA: 8.6	Edit

Experience

+ Add

TCS — Consultant	2020-23	3 bullets	Edit
------------------	---------	-----------	------

Score Against JD

AI Rewrite Bullet

Achievement Wizard

SPC Review — 2 open comments

Resolve All

PN

Priya Nair on "Experience › Bullet 2"

"Add a number here — % improvement or scale"

Apply

Dismiss

PN

Priya Nair on "Summary"

"Trim to 2 lines, too long for ATS parsers"

Apply

Dismiss

19

7. Access Control, Roles & Permission Model

7.1 Concept — Roles + Permission Sets

Two layers, kept separate so access stays composable instead of every variation needing a brand-new role — modeled on Salesforce's Profile + Permission Set pattern:

- **Role** — a user's base identity. Comes with a default bundle of Permission Sets attached. A user can hold more than one Role at once.
- **Permission Set** — a named, reusable group of actions (e.g. "JD Management"). Roles ship with default Permission Sets; Admin can also assign extra Permission Sets to an individual user directly, independent of role.
- **Custom Role** — created by cloning an existing base Role (copying its default Permission Set bundle), then adding/removing Permission Sets or individual actions within them.

7.2 Base Roles

Role	Purpose	Default Permission Sets
Admin	CDPO staff, full platform control	JD Management, Shortlist Oversight, Student Data - Full, Reports & Export, User Management, Role & Permission Management, CRM/Outreach, Audit Log View
SPC	Manages shortlist coordination logistics	Shortlist Oversight, Student Data - Full, Reports - View Only
BD	Pre-season company outreach	CRM/Outreach
Recruiter	Company-side hiring contact	JD Management (own company), Shortlisting (own JDs), Student Data - Restricted
Student	Batch candidate	Own Profile, Own Applications, Own Defaults View

All Admin-role users are equal in permissions — there is no admin hierarchy. Any Admin can approve users, create roles, and view the full audit log. A user who holds Student + SPC + BD simultaneously simply stacks all three Roles' Permission Sets.

7.3 Permission Set Catalog

Permission Set	Actions Included
JD Management	View, Create, Edit, Publish, Approve (admin gate), Delete/Archive JDs
Shortlist Oversight	View all shortlists (any company), override status, assign interview slots, send status-nudge reminders
Shortlisting (recruiter-scoped)	View own applicants, move applicant status, add private notes, download candidate packets
Student Data - Full	View all profile fields incl. contact info, bulk import/export roster, edit any student's defaults record
Student Data - Restricted	View profile fields except phone/personal email (masked until shortlist stage)

Permission Set	Actions Included
Own Profile	View/edit own profile, CV link, education details
Own Applications	Apply to JD, withdraw application, view own status timeline
Own Defaults View	View own default count/activity breakdown (read-only)
Reports & Export	View placement dashboard, export Final Placement Datasheet, export custom templates
Reports - View Only	View placement dashboard (no export)
User Management	Approve/reject accounts, deactivate/reactivate users, assign Roles and extra Permission Sets
Role & Permission Management	Create/clone/edit custom Roles and Permission Sets
CRM/Outreach	View company pipeline, log calls, run mail merge, manage outreach tasks/reminders
Audit Log View	View audit log entries (read-only for everyone, no edit/delete exists)

7.4 Field-Level Security

Independent of module-level Permission Sets — governs individual fields, not whole records. Student phone number and personal email stay masked to any "Student Data - Restricted" holder (Recruiters) until that specific candidate reaches Shortlisted status on that JD, at which point those two fields unmask for that recruiter, for that candidate only. "Student Data - Full" holders (Admin, SPC) always see everything.

7.5 User Lifecycle

Stage	Behavior
Signup	Student: SSO auto-activates. Recruiter/SPC/BD: SSO or invite signup, created in Pending status.
Approval	Any Admin approves/rejects pending Recruiter/SPC/BD accounts. Rejected users are notified.
Active	Normal access per assigned Role(s) + any extra Permission Sets.
Multiple recruiter logins	No cap per company; identical Recruiter permissions scoped to that company's own JDs; no recruiter sub-hierarchy.
Auto-expiry	Student and SPC access auto-deactivates at season end / graduation (batch-linked). Recruiter/BD accounts persist across seasons unless manually deactivated.
Manual deactivation	Available to any Admin at any time (e.g. recruiter contact leaves the company).

8. Non-Functional Notes & MVP Scope Boundaries

8.1 Explicitly Out of Scope for MVP

- Call recording/telephony integration — logged manually as notes for now.
- Multi-school/multi-tenant white-labeling — architecture should allow it later, not required now.
- Interview slot self-scheduling by students — SPC manages this manually in MVP.
- Admin hierarchy / permission restriction between Admins — all Admins are equal by design for now.

8.2 MVP Approach

Full pipeline, end-to-end, kept thin at each step — every module in Section 4 gets a working baseline in v1 rather than one module being built deep while others are absent. Depth is added module-by-module in subsequent iterations based on real season usage.

9. Open Items for Next Pass

- Full CRUD action matrix per module, beyond the illustrative actions listed in Section 7.3 — needed once build planning starts.
- Whether Permission Sets themselves should be admin-creatable (not just assignable), or remain a fixed catalog Admin composes from.
- Specific SSO provider/protocol for institute login (Section 4.9 / 7.5).
- Configurable staleness threshold default (Section 4.5) and who can change it.
- Export template ownership — who defines the accreditation-submission report format (Section 4.7).

10. Commercialization & Monetization Strategy

This section defines how Placement OS generates revenue, what it costs to acquire and serve each institute, and the incentive structure used to overcome switching friction from incumbent placement tools.

10.1 Pricing & Packaging

Core pipeline functionality (Company Onboarding, JD Management, Application Flow, Shortlisting, SPC Coordination) is priced at or near zero. The market-leading incumbent already offers this layer free to institutions, monetizing instead on the recruiter side; competing on price at this layer is not viable. Revenue is instead concentrated in two differentiated modules that incumbents do not offer at comparable depth or quality.

Offering	Price (per institute / year)	Rationale
Core Pipeline (JD, Applications, Shortlisting, SPC Coordination)	Free – ₹1L	Loss-leader; matches incumbent free tier, removes the primary adoption blocker
Resume + AI Module (standalone)	₹2L – ₹3L	Materially undercuts the known incumbent price point for a comparable module while preserving strong margin
Outreach / BD CRM Module (standalone)	₹1.5L – ₹2.5L	No strong incumbent at this depth (persona-based pitch templates, per-recipient funnel tracking); priced on value
Full Bundle (all modules)	₹4L – ₹6L	Bundling discount vs. modules purchased separately; still meaningfully below incumbent per-module pricing

10.2 Year-1 Onboarding Incentive

A fully free first year was considered and rejected as blanket policy. Two structural risks make pure-free a weak default: (1) it absorbs 100% of onboarding, LLM, and support cost with zero offsetting revenue — estimated at ₹65K–₹2.3L per institute in year one; and (2) institutional placement-committee turnover means the champion who agreed to a free pilot is frequently not the person deciding on renewal, creating a payment cliff at the exact moment institutional memory of the tool's value is weakest.

Recommended structure: waive the onboarding/setup fee (removes real adoption friction, low absolute cost) but retain a nominal commercial price from year one — e.g. ₹25K–₹50K against the ₹2L–3L sticker — with year-2 pricing locked into the year-1 agreement. This produces a signed contract and a real payment (a genuine willingness-to-pay signal, and a commitment that survives a committee handover), while still removing the bulk of first-year friction. True ₹0 pricing is reserved specifically for the IIM Raipur pilot, where the goal is a referenceable case study rather than revenue, and where direct founder relationship credibility substitutes for a commercial proof point.

10.3 Cost Structure

10.3.1 One-Time Build Cost (all 10 modules, config-driven, multi-tenant)

Module	Est. Effort (hrs)	Notes
Company Onboarding & JD Management (4.1)	40–50	Standard CRUD workflow
Eligibility Engine + Notification (4.2)	80–100	Rule-builder UI + engine, not hardcoded filters
Student Application Flow (4.3)	30–40	Straightforward
Candidate Packet & Shortlisting (4.4)	50–60	Packet compile, bulk actions, field unmask logic
Shortlist Coordination / SPC (4.5)	40–50	Dashboard, staleness flags, slot assignment
Defaults & Compliance Tracker (4.6)	60–70	Config-driven category schema per institute
Final Placement Reporting (4.7)	80–100	Report-template engine — highest-risk module
Outreach CRM (4.8)	80–120	Adapted from OutreachOS, not built fresh
Platform / Admin + RBAC (4.9, Sec.7)	120–150	Role + Permission Set engine, audit log, field-level security
Resume Maker (4.10)	100–140	Adapted from Cursivo's AI panels + SPC review flow
Infra, multi-tenant setup, SSO, QA, polish	120–150	Overhead across all modules

Total: approximately 900–1,100 hours. At blended outsourced dev cost (₹1,000–1,500/hr), this equates to ₹9L–₹16.5L cash-equivalent, one time. Solo and part-time alongside other commitments, this is roughly 8–10 months; solo and full-time, roughly 5–7 months.

10.3.2 Per-Institute Onboarding Cost

Item	Institutes 1–4 (early)	Institute 5+ (mature config tooling)
Roster import / data mapping	8–15 hrs	2–4 hrs
Defaults category schema setup	3–5 hrs	1–2 hrs
Eligibility rule configuration	2–4 hrs	1 hr
Report-template mapping (accreditation format)	8–15 hrs	2–5 hrs
SSO / institute email domain setup	3–8 hrs	3–8 hrs (institute-side variability doesn't shrink)
Branding + mailer tweaks	2–4 hrs	1–2 hrs
Onboarding / training calls	5–10 hrs	5–10 hrs
Total (hrs / cash-equivalent)	~30–55 hrs / ₹30K–₹82K	~15–35 hrs / ₹15K–₹52K

This per-institute figure only holds if the eligibility engine, defaults schema, and report-template engine are built as configurable systems in the first build pass (Section 12). Hardcoding institute #1 reverts each

subsequent institute to a near-fresh mini-build — the same trap the BD's earlier modularization plan is designed to avoid.

10.4 Margin Projection (Resume + AI Module, ₹2L–₹3L price point)

Line Item	Early institutes (1–4), Year 1	Mature institutes (5+), Renewal years
Revenue	₹2L – ₹3L	₹2L – ₹3L
Onboarding / config cost	₹30K – ₹82K (one-time)	– (already onboarded)
LLM / variable cost (per season)	₹20K – ₹1.25L	₹20K – ₹1.25L
Support / account management (time cost)	~₹15K – ₹25K equiv.	~₹10K – ₹15K equiv.
Effective margin	~40–60%	~75–80%

Founder sales time, not gross margin, is the binding constraint: closing a single institute is estimated at 20–40 hours of relationship-driven selling (calls, demos, procurement negotiation), separate from and additional to the build hours in 10.3.1.

10.5 Go-to-Market Motion

- Module-first land: sell the Resume/AI module or Outreach CRM standalone — institutes already running an incumbent core pipeline do not need to replace it to adopt a stronger module alongside it.
- Referral-driven, not marketing-driven channel: CDPO offices and placement committees coordinate informally and via inter-institute coordinator handoffs each year; paid acquisition has negligible reach into this buyer group.
- IIM Raipur pilot serves as the proof asset — founder's own placement-coordinator credibility unlocks the first 2–3 conversations, but a referenceable usage metric (e.g. outreach response-time reduction, resume review turnaround) is required before colder institutes will engage.
- Expansion beyond the pilot is module-by-module and institute-by-institute, not a platform-wide rollout — consistent with the MVP approach in Section 8.2.

11. Competitive Analysis — Superset, Reculta & Adjacent Players

Placement OS is not entering an empty market. Two established players already have meaningful footprint inside premium B-Schools, and this section positions Placement OS against them directly rather than assuming a greenfield opportunity.

11.1 Market Landscape

Superset is the segment leader, reporting 175+ colleges including IIMs and NITs, backed by Great Learning, and operating as a two-sided marketplace connecting institutions with 800+ recruiters. Its core placement pipeline is offered free or low-cost to institutions, with monetization concentrated on the recruiter side of the marketplace.

Reculta, founded in 2017 with seed backing from IIM Ahmedabad's CIIE incubator, has an existing footprint that includes IIM Ahmedabad, IIFT Delhi, IIFT Kolkata, FMS Delhi, NMIMS, XIMB, IMT, and — notably — IIM Raipur, the pilot institute named in this BRD. Reculta's product is centered on placement-process digitization and a manual, verifier-driven resume/profile-sheet workflow (points are individually flagged Verified / Pending / Rejected against supporting proof), rather than an AI-generated resume pipeline.

Smaller/adjacent players (LineupX, Oréll, Reculta-alternatives positioning) compete mainly on AI-driven interview practice and white-label flexibility rather than on placement-cell workflow depth or resume generation quality.

Critical implication for the pilot: IIM Raipur is very likely already a Reculta client, not a greenfield account. The pitch there is “replace or outperform an existing, known tool,” not “digitize a manual process from scratch” — this changes the sales conversation materially from the framing used earlier in this document and should be validated directly with the current CDPO office before pilot scoping.

11.2 Feature & Positioning Comparison

Dimension	Superset	Reculta	Placement OS
Core placement pipeline (JD, apply, shortlist)	Strong; market standard, free to institutions	Present; less widely adopted than Superset	Present; priced as loss-leader, not the differentiator
Pricing model	Free/low-cost to institution; recruiter-side monetization	Institution-paid (pricing not publicly disclosed)	Institution-paid on differentiated modules only; core near-free
Outreach / BD CRM depth (pre-season company sourcing)	Not a core focus	Not a core focus	Purpose-built: persona-based pitch templates, per-recipient funnel, call-remarks logging — direct match to the CDPO's existing JPC workflow

Dimension	Superset	Reculta	Placement OS
Defaults / compliance tracking	Generic, not institute-specific	Generic, not institute-specific	Config-driven schema per institute (Section 4.6), matching each CDPO's own category structure
Resume / CV generation	Not a core focus	Manual verifier-driven workflow (proof-based point verification), not AI-generated	AI-assisted generation + SPC inline review, adapted from Cursivo's proven engine
Report-template configurability (accreditation export)	Unclear / not a stated differentiator	Unclear / not a stated differentiator	Purpose-built report-template engine mapping to each institute's exact existing export format (Section 4.7)
Institutional footprint (as of this BRD)	175+ colleges incl. IIMs, NITs	IIM-A, IIFT, FMS, NMIMS, XIMB, IMT, IIM Raipur	Zero — pre-pilot

11.3 Where Placement OS Can Win — and Where It Can't

Genuine advantages

- AI-generated resume pipeline vs. Reculta's manual proof-verification model — materially faster turnaround for students and SPCs, and a defensible technical moat since it is adapted from a proven, already-shipped engine (Cursivo).
- Outreach/BD CRM purpose-built around the CDPO's actual JPC mailing-sheet workflow — neither Superset nor Reculta appear to compete at this depth.
- Config-driven defaults tracking and report-template engine, avoiding the generic-template problem both incumbents likely share.

Where it cannot realistically compete

- Core pipeline breadth and existing scale: Superset's 175+ college footprint and recruiter-side marketplace network effects cannot be replicated by a solo-founder build in any near-term horizon — do not position Placement OS as a full Superset replacement.
- Free-to-institution pricing on the core pipeline: matching, not beating, is the achievable goal; Placement OS's revenue must come from the modules incumbents don't offer, not from underpricing the pipeline itself.
- Brand trust / longevity: Reculta's multi-year relationship with IIM-A and other IIMs, and Superset's Great Learning backing, are trust signals a first-time pilot cannot match — the wedge-module strategy (Section 10.5) exists specifically to sidestep this rather than compete on it directly.

12. Technical Architecture & Cost Estimate

Recommended stack decisions, aligned with the founder's existing Cursivo (React/Firebase/Firestore) and OutreachOS (Next.js/Supabase) architecture, and the build-approach cost comparison used to decide how to sequence the build.

12.1 Recommended Stack

Layer	Recommendation	Rationale
Frontend / framework	Next.js	Matches OutreachOS; well suited to admin-heavy dashboards and SSR
Database	Supabase (Postgres)	Native Row-Level Security maps directly onto the Role/Permission Set model (Section 7); relational joins are required for the eligibility and reporting engines in a way Firestore's document model does not support well
Multi-tenancy	Single schema, institute_id column + RLS policies	Cheapest to build and scale to institute #10+; schema-per-tenant would recreate the slow-onboarding problem this architecture is meant to avoid
Auth / SSO	Existing Firebase Custom Token federation layer, extended	Already built; links Cursivo/OutreachOS identity — reuse rather than rebuild
Notifications	Resend	Already integrated in Cursivo
Config engine (eligibility rules, report templates)	JSON-schema-defined rules stored in Postgres, evaluated server-side	Sufficient at this scale; a heavyweight rules engine (e.g. Drools) would be over-engineering
Hosting	Vercel	Matches OutreachOS

12.2 Build Approach Comparison

Approach	What it gets you	Est. Effort	Est. Cash Cost (outsourced, ₹1,000–1,500/hr)
Single-tenant, hardcoded for IIM Raipur only	Fast pilot; proves the case study	~250–350 hrs	₹2.5L–5L
Multi-tenant, config-driven (recommended long-term)	Real “days not months” onboarding at institute #10+	~700–900 hrs	₹7L–13.5L
No-code / low-code (Bubble, Retool + Supabase)	Cheapest, fastest MVP	~150–250 hrs	₹1.5L–3.5L (likely requires a rebuild, not a

Approach	What it gets you	Est. Effort	Est. Cash Cost (outsourced, ₹1,000–1,500/hr)
			migration, once real RBAC/config is needed)
White-label / partner with a smaller player	Zero build cost, fastest to market	Near-zero	Revenue-share instead of cash cost; no differentiation or IP ownership

Recommended sequencing: build the single-tenant IIM Raipur pilot first, scoped specifically to the Resume+AI module and Outreach CRM (the two modules identified in Section 11.3 as genuine differentiators) rather than the full core pipeline. Commit to the ₹7L–13.5L multi-tenant, config-driven build only once the pilot has produced a real, referenceable usage metric — not before.

Appendix A. Decision Log

This log captures judgment calls made during BRD development that are not otherwise self-evident from reading the functional requirements. It exists to give any reader — human or AI — the reasoning behind choices that could otherwise look arbitrary or be second-guessed in a later pass.

Decision	Detail
Persona = company-type template	Each "persona" (e.g. Product/SaaS, Consulting, Core Engineering) is a reusable company-type template used to tailor outreach messaging and JD framing — not a separate sender identity. The email sender is always the logged-in user (SPC/BD owner), never a persona account.
JPC and SPC are the same role	"JPC" (Job/Placement Coordinator, per the source mailing sheet) and "SPC" (Shortlist/Student Placement Coordinator) refer to the same underlying role in PlacementOS. SPC is used as the standard term throughout the BRD; SPC additionally carries supervisory capability over other coordinators where the mailing sheet showed JPC hierarchy.
Merge status set extended	The candidate/application status set was extended beyond OutreachOS's original defaults to add two institution-specific states: "Not Interested" and "Call Back Later", both surfaced in the JPC Mailing Sheet as recurring recruiter dispositions.
JD Form Received stays manual	The "JD Form Received" flag (Section 4.1) remains a manual toggle set by the SPC/Admin. It does not auto-trigger progression into Company Onboarding — onboarding is a deliberate, separate action to allow for a review step between form receipt and pipeline entry.
Each company has a designated owner SPC	Every company in the pipeline is assigned one owner SPC for continuity of relationship and accountability, even where multiple coordinators may interact with that company during a season.
AI Writing Assist and SPC Inline Review are in MVP scope	Both the AI-assisted CV writing tools (from Cursivo's AIDiffPanel / FullRewritePanel / CoPilotPanel / AchievementWizard) and the SPC inline review/comment flow (adapted from Cursivo's ReviewInboxPanel) are included in MVP scope, not deferred to a later phase — confirmed explicitly during the Resume Maker module mapping discussion.
CV polish tools skipped for MVP	RecruiterHeatmapPanel and CVHealthPanel from Cursivo are explicitly out of MVP scope. They are nice-to-have engagement/analytics features rather than core placement workflow, and are candidates for a later phase.
Consumer-SaaS Cursivo modules excluded	Upgrade/Payment, Referral, and NaukriSync modules from Cursivo are excluded entirely from PlacementOS — these are consumer-SaaS monetization and growth features with no equivalent in an institutional placement-cell tool. Cursivo's own Admin module is likewise excluded since PlacementOS already has its own Admin section (4.9).

Appendix B. Data Validation & Migration Mapping

The Section 5 data model was defined from the source documents by inspection. This appendix re-checks it against the actual files — not summaries of them — and records what needs to change before a season-1 import can run unattended.

B.1 Students' Profile Sheet — confirmed & extended

The live file ("Students' Profile Sheet 2023-25") confirms the Student entity's core fields (roll_no, name, work-ex, up to 3 prior employers, graduation/PG/10th/12th detail, credentials). It also carries four fields not yet in the Section 5 model that recruiters and the eligibility engine will want to filter/sort on:

Source column	Add to Student as	Why it matters
S.No.	display_seq (int, non-key)	Cosmetic ordering only; not a stable identifier — do not use as a key.
Section	section (string)	Batch is split into sections; SPCs and Admin filter rosters by section.
Age	age (int)	Occasionally an eligibility criterion for public-sector recruiters.
Gender	gender (string)	Some recruiters filter by diversity hiring criteria; must stay masked pre-shortlist like phone/email (FR-3 area).

B.2 Final Placement Datasheet — structure risk

This file is not a normalized table — it is a pivoted, human-formatted layout: one 4-column block per company (Date Floated / Date of Process / Profiles Offered / a results column), with student-level placement rows beneath. That's very different from the "wide-table format (company ↔ student ↔ package)" description in Section 4.7, and it directly affects FR-7.2 (one-click export in this format) and the Admin import step in Section 3.4 Step 1. Concretely, inspecting the live file surfaced:

- An '#REF!' formula error sitting in the Average Package summary cell — any auto-generated export in this format needs its own formulas, not a copy of the source sheet's.
- Date fields containing 'TBD' as a literal string instead of being blank — a naive date parser will throw on these; the import needs an explicit TBD/unknown state, not a parse error.
- Some 'date' cells hold multiple comma-separated dates as text (e.g. multiple interview dates in one cell) rather than one value — these map to round_history[] on Application, not a single DefaultRecord-style scalar.
- 247 companies and 337 placed / 42 yet-to-be-placed students in the 2024-25 season is the real scale the pilot's eligible-list and shortlist queries need to perform well against (see Appendix C).

Recommendation: treat season-1 migration of this file as a reviewed, semi-manual import (load into staging, show Admin a diff/validation screen, then commit) rather than a one-click CSV import — the source data itself needs a human pass, not just a parser.

B.3 JPC Mailing Sheet — richer than modeled

The live sheet (18 tabs, one per JPC, ~20 columns each) confirms the OutreachActivity/CompanyContact shape but is more granular than Section 5 currently reflects:

Source column	Model implication
JPC / JPC Line / JPCTitle	These are the JPC's own name/phone/title on each row, not just a logged_by user_id — needed so imported historical rows can be attributed even before those JPCs have system accounts.
Call Remarks / SPC Remarks / JPC Remark	Three distinct free-text fields, not the two (call_remarks, spc_remarks) currently in OutreachActivity — add jpc_remark.
SPC Name / SPC Mobile No	Supervisor contact captured per-row in the sheet; maps to the supervisor_user_id relationship on Company, but needs a text fallback for pre-migration rows where that SPC has no user record yet.
Merge status	Stored as coded tokens (EMAIL_SENT, EMAIL_CLICKED, ...), not the human-readable labels used in Section 3.5/4.8 prose — needs an explicit enum mapping table (below), not a direct string copy.
Sheet value	PlacementOS enum
EMAIL_SENT	Email Sent
EMAIL_CLICKED	Clicked
(observed, not yet in sample)	Opened / Responded / Not Interested / Call Back Later / Bounced — confirm full token set against all 18 tabs before building the mapping table, the two-row sample above doesn't cover it.

Appendix C. Success Metrics & KPIs

Metrics for the single-tenant IIM Raipur pilot (Section 12.2). Baselines below come from the actual 2024-25 Final Placement Datasheet, so the pilot has a real before/after to report rather than an estimate.

Metric	2024-25 baseline (manual process)	MVP target
Companies processed per season	247	No regression — same or higher at equal CDPO staff effort
Students placed / yet-to-be-placed	337 placed / 42 yet-to-be-placed	Track live throughout the season instead of only at reporting checkpoints
Time from JD receipt to notification going out	Manual compilation + mailer send (hours, ad hoc)	< 30 minutes from Publish to notification delivery (FR-2.4/2.6)
Candidate packet compilation	Manually assembled per company	Auto-compiled at application time (FR-4.1) — zero manual compilation
Stale pipeline detection	Manual tracking in per-JPC sheets/notebooks	System-flagged automatically within the configured staleness window (FR-5.4)
Final Datasheet turnaround	Manual, end-of-season	One-click export (FR-7.2), same format, on demand mid-season too

Adoption metrics to watch during the pilot: % of recruiters who self-register vs. need Admin to onboard them manually; % of JPCs who log outreach in-system vs. reverting to personal notes; SPC daily-active-use rate during peak shortlisting weeks.

Appendix D. Risks & Assumptions

Risk / assumption	Impact if wrong	Mitigation
Historical data (Profile Sheet, Datasheet, JPC sheets) imports cleanly	Season-1 launch delayed by manual data cleanup (see Appendix B.2)	Budget an explicit staging + Admin-review step in the import flow, not a one-shot script
CDPO office and SPCs will adopt a new tool mid-process	Falls back to spreadsheets/email, pilot produces no usable metric	Launch ahead of a season start (not mid-season); keep the student-facing mailer format identical (FR-2.4) to lower the change felt by students
Institute IT will support/allow an SSO or work-email domain check for recruiter signup	FR-1.1 self-registration gate weakens to fully manual approval	Confirm domain-based approval is acceptable without full SSO before build; Admin manual approval remains the fallback either way
Usage is seasonal (spiky around JD floating and shortlist weeks)	Under-provisioned infra during peak weeks; idle cost off-season	Pick a hosting approach in Section 12.1 that scales down between seasons, not fixed-capacity
Single owner (Chittaranjan) builds and supports this alongside other IITIIM Careers products	Support/bug-fix latency during a live placement season is high-stakes for students' offers	Scope the pilot to the two differentiator modules first (Resume+AI, Outreach CRM) per Section 12.2, not the full pipeline, to keep surface area supportable

Appendix E. Security, Privacy & Compliance

Student records (phone, personal email, marks, defaults status, employment history) are personal data under India's Digital Personal Data Protection Act, 2023 (DPDP Act). This section adds the compliance framing that Sections 4 and 7's field-masking rules (e.g. FR applicant contact fields staying masked until shortlist) already imply but don't name.

- Data fiduciary vs. processor: the institute (CDPO) is the data fiduciary for student records; Placement OS operates as processor. A data-processing clause in the institute agreement should say so explicitly, alongside who's on the hook for a breach notification.
- Purpose limitation: student data collected for placement should not be reused for unrelated purposes (e.g. marketing) without separate consent — relevant if the platform later cross-sells Cursivo/OutreachOS to the same student base.
- Consent & notice at signup: student accounts are created via bulk roster import (Section 3.4 Step 1), not self-signup, so consent capture needs to happen at the institute level (e.g. as part of program enrollment) rather than assumed from platform use.
- Field-level masking (Section 7.4) already covers phone/personal-email pre-shortlist; extend the same treatment to the newly identified 'gender' field (Appendix B.1) and to defaults/compliance detail, which is sensitive in a different way (academic standing).
- Retention & deletion: Section 7.5 covers account expiry at season end/graduation, but the BRD doesn't yet state a deletion timeline for placed students' full profile+CV packets after they've served their purpose — worth setting an explicit retention period (e.g. N years post-graduation for accreditation records) rather than indefinite retention by default.
- Recruiter-side data: company HR contacts' emails/phone numbers are also personal data; the same purpose-limitation logic applies if outreach lists are ever reused across institutes on a multi-tenant build (Section 12.2's later phase).

Appendix F. MVP Acceptance Criteria / Definition of Done

Per-module checklist for calling the single-tenant IIM Raipur pilot (Section 12.2) done. Each line should be demonstrable against real data, not synthetic test data, before sign-off.

Company Onboarding & JD Management

- A recruiter can self-register, get Admin-approved, and publish a JD without CDPO staff touching the JD content.
- Eligible-student count preview (FR-1.4) matches a manual count for at least 3 real JDs from the 2024-25 roster.

Eligibility Engine & Notification

- A JD scoped to 'unplaced only' correctly excludes all 337 already-placed 2024-25 students when replayed against last season's data.
- Notification email renders in the existing CDPO mailer format (FR-2.4), verified side-by-side against a real sample circular.

Student Application & Candidate Packet

- One-click apply attaches the correct Profile Sheet + latest CV with zero re-entry (FR-3.1/4.1) for a student with all three prior-employer slots filled.
- Applicant list sort/filter by CGPA and work-ex returns the same ordering as a manual spreadsheet sort on the same data.

Defaults & Compliance Tracker

- Bulk CSV import of defaults data reproduces the Total Defaults column from the source Defaults Tracker exactly, for the full batch.
- A student blocked by defaults sees the specific blocking reason (FR-3.4), not a generic 'ineligible' message.

Final Placement Reporting

- Exported Final Placement Datasheet reconciles to the real 2024-25 numbers (247 companies, 337 placed) when run against migrated historical data.
- The #REF! and TBD issues found in Appendix B.2 do not appear in any system-generated export.

Pre-Season Outreach CRM

- A JPC can log a call and see it reflected in the Senior SPC's supervision view within the same session.
- All merge_status values observed across the 18 real JPC tabs map cleanly to the system enum (Appendix B.3) with none falling into an 'unmapped' bucket.

Appendix G. Glossary

Term	Meaning
CDPO	Career Development & Placement Office — the institute department running placements.
SPC	Student Placement Coordinator; student volunteer running part of the process, incl. outreach.
JPC	Job/Placement Coordinator — the CDPO's original term for the same SPC-role outreach function (Decision Log, Appendix A).
GL	Guest Lecture — one of the tracked activity types feeding the Defaults Tracker.
Defaults	Attendance/compliance shortfalls (GLs, seminars, summits, process defaults) that can block placement eligibility.
Merge status	The outreach-email delivery/response state for a company contact (Sent/Opened/Clicked/Responded/etc.).
CTC	Cost to Company — total compensation, usually broken into fixed/variable components on a JD.
Unplaced-only JD	A JD visible only to students who have not yet accepted an offer this season.
Season	One placement cycle for a batch, e.g. PGP 2024-26; the unit Section 4.7's batch-over-batch comparison operates on.